

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from matplotlib.ticker import FormatStrFormatter

'''
Goal of the project:
Creating predictive models to see whether a customer will churn(terminate contra
given a dataset of customer information
'''
```

```
In [2]: df = pd.read_csv("/Users/pacosun/Downloads/Telco-Customer-Churn.csv")
```

```
In [3]: df.head()

'''
Each customer comes with 21 different features which will help us
create predictive model later in the project

Will not be using every single one; some trash features will be dropped
'''
```

```
Out[3]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No

5 rows × 21 columns

```
In [4]: # Quick overview, confirming there's no missing values (NaN) using .isna() to co

df.isna().sum()
```

```
Out[4]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents   0
tenure       0
PhoneService  0
MultipleLines  0
InternetService  0
```

```

OnlineSecurity      0
OnlineBackup        0
DeviceProtection    0
TechSupport         0
StreamingTV         0
StreamingMovies     0
Contract            0
PaperlessBilling    0
PaymentMethod       0
MonthlyCharges      0
TotalCharges        0
Churn               0
dtype: int64

```

In [5]:


```

# Displaying the balance with CountPlot

plt.figure(figsize = (6,5))
sns.countplot(data = df, x = "Churn")
plt.show()

'''
Overall, the majority did not churn
However more specific analysis needed
'''

```



In [6]:

```


# Creating BoxPlots for distribution of TotalCharges per contract type
# Adding hue coloring based on Churn class

plt.figure(figsize = (10,6))
sns.boxplot(data = df, x = "TotalCharges", y = "Contract",
            hue = "Churn")
plt.legend(loc = (1.01,.5))
plt.show()

'''
Month-to-month is hard to predict because customers come expecting to cancel the

```

Two-year does the telling: those who did churn were charged higher (also higher ...



```
In [7]: # Selecting features with less unique instances
# customerID, tenure will not be selected to convert to dummy values due to no o
# -> highly unique features
```

```
corr = pd.get_dummies(df[['gender', 'SeniorCitizen', 'Partner',
                          'Dependents', 'PhoneService', 'MultipleLines',
                          'InternetService', 'OnlineSecurity', 'OnlineBackup',
                          'DeviceProtection', 'TechSupport', 'StreamingTV',
                          'StreamingMovies', 'Contract', 'PaperlessBilling',
                          'PaymentMethod', 'Churn']]).corr()
```

```
In [8]: # Creating correlation for features selected above

corr["Churn_Yes"].sort_values().iloc[1:-1].head()

...
Two-year has a correlation of -.301, meaning it's not likely for someone who's o
contract to churn
...
```

```
Out[8]: Contract_Two year                -0.301552
StreamingMovies_No internet service    -0.227578
StreamingTV_No internet service        -0.227578
TechSupport_No internet service        -0.227578
DeviceProtection_No internet service   -0.227578
Name: Churn_Yes, dtype: float64
```

```
In [9]: # Visualization of correlations

plt.figure(figsize = (15,10), dpi = 200)
sns.barplot(x = corr["Churn_Yes"].sort_values().iloc[1:-1].index,
            y = corr["Churn_Yes"].sort_values().iloc[1:-1].values)
plt.title("Correlation to Churn_Yes")
```


```
plt.xticks(rotation = 90)
plt.show()
```

```
'''
```

Again, it's unlikely for one to churn if has a two-year contract
 -> Two-year contract on the far left with a negative correlation

On the contrary one would be a lot more likely to churn if on a monthly contract
 -> Month-to-month on the far right with a positive correlation


```
'''
```



In [10]:

```
# Distributing based on tenure -> (Amount of time one stays as a customer)
# Visualizing the distribution of tenure using histogram
```

```
plt.figure(figsize = (10,7))
sns.histplot(data = df, x = "tenure", bins = 40)
plt.show()
```



```
In [11]: # Grouping based on Tenure length (1 Month, 2 Months, ... n Months)
# Treating each length as a single group/cohort
# And then calculate the "Churn Rate" according to tenure

noChurn = df.groupby(["Churn", "tenure"]).count().transpose()["No"]
yesChurn = df.groupby(["Churn", "tenure"]).count().transpose()["Yes"]
```

```
In [12]: churnRate = yesChurn / (noChurn + yesChurn) * 100
```

```
In [13]: churnRate.transpose()["customerID"]

...


List below shows there's a negative correlation: LONGER the tenure, LESS likely
-> e.g. One-month tenure would churn 61.99% percent of the time
...
```

```
Out[13]: tenure
1      61.990212
2      51.680672
3      47.000000
4      47.159091
5      48.120301
...
68      9.000000
69      8.421053
70      9.243697
71      3.529412
72      1.657459
Name: customerID, Length: 72, dtype: float64
```

In [14]:

```
# Visualizing the correlation

plt.figure(figsize = (10,5),dpi = 100)
churnRate.iloc[0].plot()
plt.ylabel("Churn Rate")
plt.show()
```



In [15]:

```
# Broadening the cohorts
# Separating into 4 groups
# 0 - 12, 12 - 24, 24 - 48, Over 48

def cohorts(n):
    if n < 13:
        return "0 to 12 Months"
    elif n < 25:
        return "12 to 24 Months"
    elif n < 49:
        return "24 to 48 Months"
    else:
        return "Over 48 Months"
```

In [16]:

```
# Adding a new column -> Tenure Group

df["Tenure Group"] = df["tenure"].apply(cohorts)
```

In [17]:

```
df.head(11)[["tenure", "Tenure Group"]]
```

Out[17]:


	tenure	Tenure Group
0	1	0 to 12 Months
1	34	24 to 48 Months
2	2	0 to 12 Months

	tenure	Tenure Group
3	45	24 to 48 Months
4	2	0 to 12 Months
5	8	0 to 12 Months
6	22	12 to 24 Months
7	10	0 to 12 Months
8	28	24 to 48 Months
9	62	Over 48 Months
10	13	12 to 24 Months

In [18]:

```
# Visualizing the relationship between Monthly and TotalCharges
# Colored by Tenure Group

plt.figure(figsize = (10,5), dpi = 100)
sns.scatterplot(data = df, x = "MonthlyCharges",
                y = "TotalCharges", hue = "Tenure Group", alpha = .5)
plt.show()
```




In [19]:

```
# Creating a grid of CountPlots
# Showing counts per Tenure Group, separated by contract type
# Colored by Churn

plt.figure(figsize = (10,8), dpi = 100)
sns.catplot(data = df, x = "Tenure Group",
            hue = "Churn", col = "Contract", kind = "count")
plt.show()

'''
Here we are seeing a significant drop in churn rate as the contract length(type)
'''
```

<Figure size 1000x800 with 0 Axes>



```
In [20]: # Now deploying tree based methods for predictive modeling
# DecisionTree, RandomForest, BoostedTrees(Ada & Gradient)
```

```
In [21]: # Separating data into X and y label
# Creating dummy values

X = df.drop(["Churn", "customerID"], axis = 1)
X = pd.get_dummies(X, drop_first = True)
y = df["Churn"]
```

```
In [22]: # Performing a train-test split
# 10% for testing, randomState = 101

from sklearn.model_selection import train_test_split
```

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                            test_size = .1,
                                                            random_state = 101)
```

```
In [24]: # 1. Train DecisionTree model

# 2. Evaluate performance metrics
# -> Classification report, confusion matrix

# 3. Calculate feature importance
# 4. Plot the tree

from sklearn.tree import DecisionTreeClassifier
```

```
In [25]: tree = DecisionTreeClassifier(max_depth = 6)
```

```
In [26]: tree.fit(X_train, y_train)
```

```
Out[26]: DecisionTreeClassifier(max_depth=6)
```

```
In [27]:
```



```
prediction = tree.predict(X_test)
```

```
In [28]: from sklearn.metrics import accuracy_score, plot_confusion_matrix, classification_report
```

```
In [29]: # Comparison between true values and predictions

print(classification_report(y_test, prediction))


'''
The model performed better when it comes to predicting those who did NOT churn c
-> (.87 vs .55)
'''
```

	precision	recall	f1-score	support
No	0.87	0.89	0.88	557
Yes	0.55	0.49	0.52	147
accuracy			0.81	704
macro avg	0.71	0.69	0.70	704
weighted avg	0.80	0.81	0.81	704

```
In [30]: import warnings
warnings.filterwarnings('ignore')
```

```
In [31]: plot_confusion_matrix(tree,X_test,y_test)
plt.show()

'''
Most important to note & reduce:
The model predicted that 75 people were not going to churn, whereas in reality t
'''
```



```
In [32]: # Calculating feature importances

features = pd.DataFrame(data = tree.feature_importances_,
```

```
index = X.columns,
columns = ["Feature Importance"]).sort_values("Feature I
```


In [33]:

```
plt.figure(figsize = (15,8),dpi=150)
sns.barplot(data = features.sort_values('Feature Importance'),
            x=features.sort_values('Feature Importance').index,
            y='Feature Importance')
plt.xticks(rotation=90)
plt.title("Feature Importance DecisionTree")
plt.show()

'''
Here we can see which features are 0, meaning that they are not at all important

*Codes to make the graph look neater:
-> features = feature[feature["Feature Importance"] > 0]

This line, if run, gets rid of features with value = 0
'''
```



In [34]:

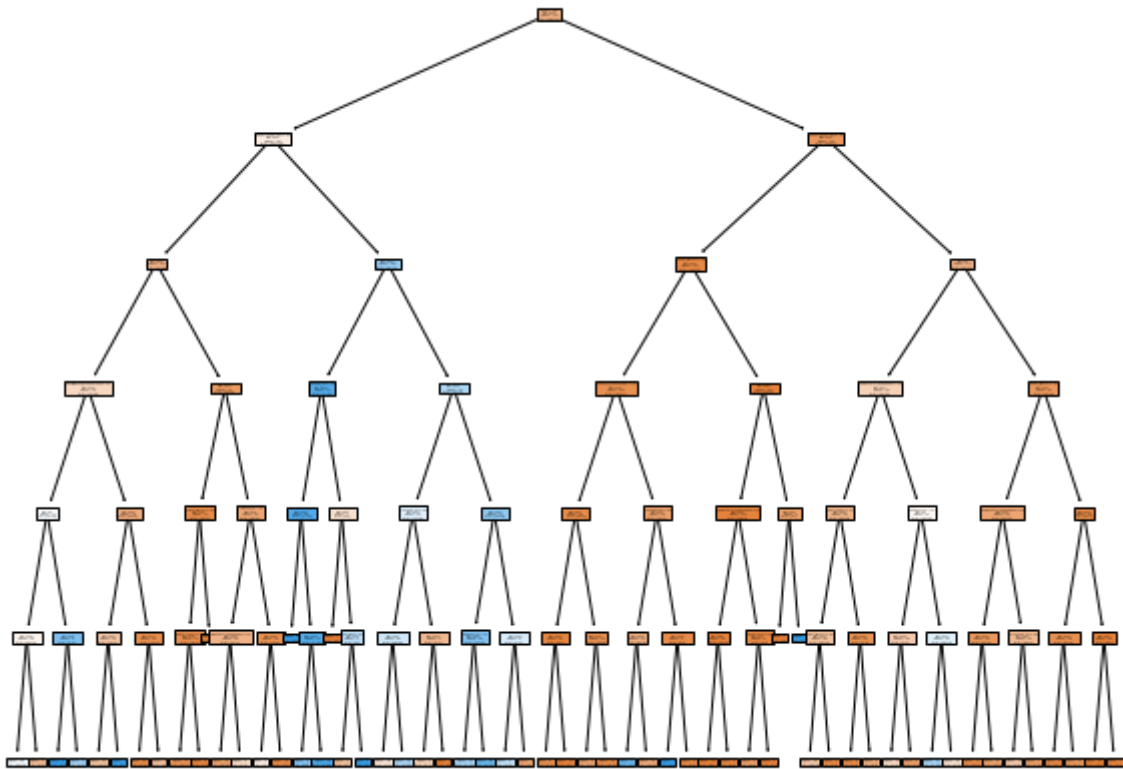
```
from sklearn.tree import plot_tree
```

In [36]:

```
plt.figure(figsize = (10,8))
plot_tree(tree, filled = True, feature_names = X.columns);

'''
```

Not necessarily a practical model since Jupyter Notebook does not allow zooming
'''



```
In [37]: # Creating a RandomForest model, and  
# a classification report, and  
# a confusion matrix from predicted values  
  
from sklearn.ensemble import RandomForestClassifier
```

```
In [51]: random = RandomForestClassifier(n_estimators = 100,  
                                       max_depth = 6)
```

```
In [52]: random.fit(X_train, y_train)
```

```
Out[52]: RandomForestClassifier(max_depth=6)
```

```
In [53]: predictions = random.predict(X_test)
```

```
In [58]: print(classification_report(y_test, predictions))  
  
'''  
Slightly worse than DecisionTree when using default value -> maxDepth = None  
-> accuracy = .80  
  
Model became better when maxDepth = 6 (same as DecisionTree)
```


```
-> accuracy = .82
'''
```

	precision	recall	f1-score	support
No	0.86	0.93	0.89	557
Yes	0.61	0.43	0.50	147
accuracy			0.82	704
macro avg	0.73	0.68	0.70	704
weighted avg	0.81	0.82	0.81	704

In [59]:

```
plot_confusion_matrix(random,X_test,y_test)
plt.show()

'''
Overall better performance than DecisionTree despite predicting that 84 customer
whereas they did have churned.
'''
```



In [44]:

```
# Creating a model using AdaBoost & GradientBoosting,
# reporting back the classification report, and
# plotting a confusion matrix

from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier
```

In [45]:

```
ada = AdaBoostClassifier()
```

In [46]:

```
ada.fit(X_train, y_train)
```

Out[46]:

```
AdaBoostClassifier()
```

In [60]:

```
adaPreds = ada.predict(X_test)
```

In [61]:

```
gboost = GradientBoostingClassifier()
```

```
In [62]: gboost.fit(X_train, y_train)
```

```
Out[62]: GradientBoostingClassifier()
```

```
In [64]: gboostPreds = gboost.predict(X_test)
```

```
In [65]: print(classification_report(y_test, adaPreds))
```


	precision	recall	f1-score	support
No	0.88	0.90	0.89	557
Yes	0.60	0.54	0.57	147
accuracy			0.83	704
macro avg	0.74	0.72	0.73	704
weighted avg	0.82	0.83	0.83	704

```
In [66]: print(classification_report(y_test, gboostPreds))
```

	precision	recall	f1-score	support
No	0.87	0.90	0.89	557
Yes	0.57	0.50	0.53	147
accuracy			0.82	704
macro avg	0.72	0.70	0.71	704
weighted avg	0.81	0.82	0.81	704

```
In [ ]: ...  
Despite both models being similarly accurate, AdaBoost recall rate is quite high  
-> .54 vs .50  
  
This statline indicates that AdaBoost identifies those who were going to churn m  
  
Therefore confusion matrix will be based on AdaBoost  
...
```

```
In [67]: plot_confusion_matrix(ada, X_test, y_test)  
plt.show()
```



In [50]:

```
...  
To sum up, we got the best performance from AdaBoostClassifier  
  
Potential future improvement for this project:  
  
1. Perform GridSearching for optimal hyperparameters since none of the models si  
those who were going to churn.  
  
2. Focus more on adjusting n_estimators such as doubling up  
...
```